Process Synchronization

Background

- Processes can execute concurrently
  - May be interrupted at any time, partially completing execution

- Concurrent access to shared data may result in data inconsistency

- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

- Illustration of the problem:
  Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer `counter` that keeps track of the number of full buffers. Initially, `counter` is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.

Producer

```
while (true) {
        /* produce an item in next produced */

        while (counter == BUFFER_SIZE) ;
                /* do nothing */
        buffer[in] = next_produced;
        in = (in + 1) % BUFFER_SIZE;
        counter++;
}
```

Consumer

```
while (true) {
        while (counter == 0)
                ; /* do nothing */
        next_consumed = buffer[out];
        out = (out + 1) % BUFFER_SIZE;
         counter--;
        /* consume the item in next consumed */
}
```

Race Condition

- **counter++** could be implemented as

  ```
  register1 = counter
  register1 = register1 + 1
  counter = register1
  ```

- **counter--** could be implemented as

  ```
  register2 = counter
  register2 = register2 - 1
  counter = register2
  ```

- Consider this execution interleaving with "count = 5" initially:

S0: producer execute **register1 = counter**	{register1 = 5}
S1: producer execute **register1 = register1 + 1**	{register1 = 6}
S2: consumer execute **register2 = counter**	{register2 = 5}
S3: consumer execute **register2 = register2 – 1**	{register2 = 4}
S4: producer execute **counter = register1**	{counter = 6 }
S5: consumer execute **counter = register2**	{counter = 4}

Critical Section Problem

- Consider system of *n* processes $\{p_0, p_1, \ldots p_{n-1}\}$

- Each process has **critical section** segment of code
  - Process may be changing common variables, updating table, writing file, etc
  - When one process in critical section, no other may be in its critical section

- ***Critical section problem*** is to design protocol to solve this

- Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**

Critical Section

- General structure of process P_i

```
do {

    [entry section]

        critical section

    [exit section]

        remainder section

} while (true);
```

```
do {

        while (turn == j);

                critical section

    turn = j;

                remainder section

} while (true);
```

Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections

2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely

3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted

   - Assume that each process executes at a nonzero speed
   - No assumption concerning **relative speed** of the n processes

Critical-Section Handling in OS

Two approaches depending on if kernel is preemptive or non-preemptive

- **Preemptive** – allows preemption of process when running in kernel mode

- **Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU

  - Essentially free of race conditions in kernel mode

Peterson's Solution

- Good algorithmic description of solving the problem

- Two process solution

- Assume that the **load** and **store** machine-language instructions are atomic; that is, cannot be interrupted

- The two processes share two variables:
  - `int turn;`
  - `Boolean flag[2]`

- The variable `turn` indicates whose turn it is to enter the critical section

- The `flag` array is used to indicate if a process is ready to enter the critical section. `flag[i] = ` *true* implies that process P_i is ready!

Algorithm for Process P$_i$

```
do {

        flag[i] = true;

        turn = j;

        while (flag[j] && turn = = j);

                critical section

        flag[i] = false;

                remainder section

} while (true);
```

Peterson's Solution (Cont.)

- Provable that the three CS requirement are met:
  1. Mutual exclusion is preserved

     P_i enters CS only if:

     either `flag[j] = false` or `turn = i`
  2. Progress requirement is satisfied
  3. Bounded-waiting requirement is met

Synchronization Hardware

- Many systems provide hardware support for implementing the critical section code.

- All solutions below based on idea of **locking**

  - Protecting critical regions via locks

- Uniprocessors – could disable interrupts

  - Currently running code would execute without preemption

  - Generally too inefficient on multiprocessor systems

    - Operating systems using this not broadly scalable

- Modern machines provide special atomic hardware instructions

    - **Atomic** = non-interruptible

  - Either test memory word and set value

  - Or swap contents of two memory words

Solution to Critical-section Problem Using Locks

```
do {

        acquire lock

        critical section

        release lock

        remainder section

} while (TRUE);
```

test_and_set Instruction

Definition:

```
boolean test_and_set (boolean *target)
  {
       boolean rv = *target;
       *target = TRUE;
       return rv:
  }
```

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to "TRUE".

Solution using test_and_set()

- Shared Boolean variable lock, initialized to FALSE
- Solution:

```
do {
        while (test_and_set(&lock))
        ; /* do nothing */
            /* critical section */
    lock = false;
            /* remainder section */
} while (true);
```

compare_and_swap Instruction

Definition:

```
int compare _and_swap(int *value, int expected, int new_value) {
    int temp = *value;

    if (*value == expected)
        *value = new_value;

    return temp;
}
```

1. Executed atomically
2. Returns the original value of passed parameter "value"
3. Set the variable "value" the value of the passed parameter "new_value" but only if "value" =="expected". That is, the swap takes place only under this condition.

Solution using compare_and_swap

- Shared integer "lock" initialized to 0;
- Solution:

```
do {
    while (compare_and_swap(&lock, 0, 1) != 0)
     ; /* do nothing */
   /* critical section */
 lock = 0;
   /* remainder section */
} while (true);
```

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    waiting[i] = false;
    /* critical section */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false;
    else
        waiting[j] = false;
    /* remainder section */
} while (true);
```

Mutex Locks

- Previous solutions are complicated and generally inaccessible to application programmers

- OS designers build software tools to solve critical section problem

- Simplest is mutex lock

- Protect a critical section by first **`acquire()`** a lock then **`release()`** the lock

  - Boolean variable indicating if lock is available or not

- Calls to **`acquire()`** and **`release()`** must be atomic

  - Usually implemented via hardware atomic instructions

- But this solution requires **busy waiting**

  - This lock therefore called a **spinlock**

acquire() and release()

- ```
  acquire() {
      while (!available)
          ; /* busy wait */
      available = false;
  }
  ```
- ```
  release() {
      available = true;
  }
  ```
- ```
  do {
      acquire lock
          critical section
      release lock
          remainder section
  } while (true);
  ```

Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.

- Semaphore **S** – integer variable

- Can only be accessed via two indivisible (atomic) operations
  - **wait()** and **signal()**
    - Originally called **P()** and **V()**

- Definition of the **wait() operation**

```
wait(S) {
    while (S <= 0)
        ; // busy wait
    S--;
}
```

- Definition of the **signal() operation**

```
signal(S) {
    S++;
}
```

Semaphore Usage

- **Counting semaphore** – integer value can range over an unrestricted domain

- **Binary semaphore** – integer value can range only between 0 and 1
    - Same as a **mutex lock**

- Can solve various synchronization problems

- Consider P_1 and P_2 that require S_1 to happen before S_2

    Create a semaphore "**synch**" initialized to 0

    ```
    P1:
        S₁;
        signal(synch);
    P2:
        wait(synch);
        S₂;
    ```

- Can implement a counting semaphore **S** as a binary semaphore

Semaphore Implementation

- Must guarantee that no two processes can execute the `wait()` and `signal()` on the same semaphore at the same time

- Thus, the implementation becomes the critical section problem where the `wait` and `signal` code are placed in the critical section

  - Could now have **busy waiting** in critical section implementation

    ‣ But implementation code is short

    ‣ Little busy waiting if critical section rarely occupied

- Note that applications may spend lots of time in critical sections and therefore this is not a good solution

Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue

- Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record in the list

- Two operations:
  - **block** – place the process invoking the operation on the appropriate waiting queue
  - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue

- ```
  typedef struct{
  int value;
  struct process *list;
  } semaphore;
  ```

```
wait(semaphore *S) {

    S->value--;

    if (S->value < 0) {
        add this process to S->list;

        block();

    }

}


signal(semaphore *S) {

    S->value++;

    if (S->value <= 0) {
        remove a process P from S->list;

        wakeup(P);

    }

}
```

Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

- Let *S* and *Q* be two semaphores initialized to 1

P_0	P_1
`wait(S);`	`wait(Q);`
`wait(Q);`	`wait(S);`
…	...
`signal(S);`	`signal(Q);`
`signal(Q);`	`signal(S);`

- **Starvation** – **indefinite blocking**
  - A process may never be removed from the semaphore queue in which it is suspended

- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
  - Solved via **priority-inheritance protocol**

Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes

    - Bounded-Buffer Problem

    - Readers and Writers Problem

    - Dining-Philosophers Problem

Bounded-Buffer Problem

- ***n*** buffers, each can hold one item

- Semaphore `mutex` initialized to the value 1

- Semaphore `full` initialized to the value 0

- Semaphore `empty` initialized to the value n

- The structure of the producer process

```
do {
    ...
     /* produce an item in next_produced */
    ...
   wait(empty);
   wait(mutex);

     ...
     /* add next produced to the buffer */
     ...
   signal(mutex);
   signal(full);
} while (true);
```

- The structure of the consumer process

```
Do {
    wait(full);
    wait(mutex);

        ...
        /* remove an item from buffer to next_consumed */

        ...
    signal(mutex);
    signal(empty);

        ...
        /* consume the item in next consumed */

        ...
} while (true);
```

Readers-Writers Problem

- A data set is shared among a number of concurrent processes
  - Readers – only read the data set; they do *not* perform any updates
  - Writers   – can both read and write
- Problem – allow multiple readers to read at the same time
  - Only one single writer can access the shared data at the same time
- Several variations of how readers and writers are considered – all involve some form of priorities
- Shared Data
  - Data set
  - Semaphore `rw_mutex` initialized to 1
  - Semaphore `mutex` initialized to 1
  - Integer `read_count` initialized to 0

- The structure of a writer process

```
do {
     wait(rw_mutex);

       ...
     /* writing is performed */

       ...
    signal(rw_mutex);
} while (true);
```

Readers-Writers Problem (Cont.)

- The structure of a reader process

```
do {
        wait(mutex);
        read_count++;
        if (read_count == 1)
         wait(rw_mutex);
     signal(mutex);

          ...
        /* reading is performed */

          ...
     wait(mutex);
        read count--;
        if (read_count == 0)
     signal(rw_mutex);

     signal(mutex);
} while (true);
```

Readers-Writers Problem Variations

- *First* variation – no reader kept waiting unless writer has permission to use shared object

- *Second* variation – once writer is ready, it performs the write ASAP

- Both may have starvation leading to even more variations

- Problem is solved on some systems by kernel providing reader-writer locks

Dining-Philosophers Problem

- Philosophers spend their lives alternating thinking and eating

- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl

  - Need both to eat, then release both when done

- In the case of 5 philosophers

  - Shared data

    ‣ Bowl of rice (data set)

    ‣ Semaphore chopstick [5] initialized to 1

Dining-Philosophers Problem Algorithm

- The structure of Philosopher *i*:

```
do {
    wait (chopstick[i] );
    wait (chopStick[ (i + 1) % 5] );

            //  eat

    signal (chopstick[i] );
    signal (chopstick[ (i + 1) % 5] );

            //  think

} while (TRUE);
```

- What is the problem with this algorithm?

Dining-Philosophers Problem Algorithm (Cont.)

- Deadlock handling

  - Allow at most 4 philosophers to be sitting simultaneously at the table.

  - Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section.

  - Use an asymmetric solution -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick.

Problems with Semaphores

- Incorrect use of semaphore operations:

  - signal (mutex) …. wait (mutex)

  - wait (mutex) … wait (mutex)

  - Omitting of wait (mutex) or signal (mutex) (or both)

- Deadlock and starvation are possible.

Monitors

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- *Abstract data type*, internal variables only accessible by code within the procedure
- Only one process may be active within the monitor at a time
- But not powerful enough to model some synchronization schemes

```
monitor monitor-name
{
  // shared variable declarations
  procedure P1 (…) { …. }

  procedure Pn (…) {……}

    Initialization code (…) { … }
  }
}
```

Schematic view of a Monitor

Condition Variables

- **`condition x, y;`**
- Two operations are allowed on a condition variable:
  - **`x.wait()`** – a process that invokes the operation is suspended until **`x.signal()`**
  - **`x.signal()`** – resumes one of processes (if any) that invoked **`x.wait()`**
    - ‣ If no **`x.wait()`** on the variable, then it has no effect on the variable

Condition Variables Choices

- If process P invokes `x.signal(),` and process Q is suspended in `x.wait()`, what should happen next?

  - Both Q and P cannot execute in paralel. If Q is resumed, then P must wait

- Options include

  - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition

  - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition

  - Both have pros and cons – language implementer can decide

  - Monitors implemented in Concurrent Pascal compromise

    ‣ P executing signal immediately leaves the monitor, Q is resumed

  - Implemented in other languages including Mesa, C#, Java

Monitor Solution to Dining Philosophers

```
monitor DiningPhilosophers
{
    enum { THINKING; HUNGRY, EATING) state [5] ;
    condition self [5];

    void pickup (int i) {
            state[i] = HUNGRY;
            test(i);
            if (state[i] != EATING) self[i].wait;
    }


    void putdown (int i) {
            state[i] = THINKING;
                        // test left and right neighbors
             test((i + 4) % 5);
             test((i + 1) % 5);
    }
```

```
void test (int i) {
        if ((state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {
              state[i] = EATING ;
            self[i].signal () ;
        }
}

    initialization_code() {
      for (int i = 0; i < 5; i++)
      state[i] = THINKING;
     }
}
```

Solution to Dining Philosophers (Cont.)

- Each philosopher *i* invokes the operations **pickup()** and **putdown()** in the following sequence:

    DiningPhilosophers.pickup(i);

    EAT

    DiningPhilosophers.putdown(i);

- No deadlock, but starvation is possible