

Buffer Cache

Tanaji Patil

Department of Computer Science and Engineering



- 1 Introduction
- 2 The Buffer Cache
- 3 Retrieval of a Buffer
- 4 Algorithms

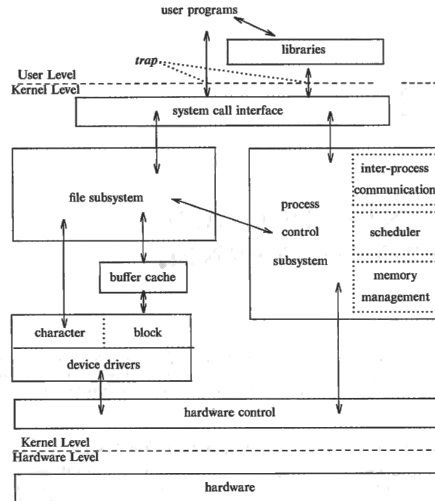
BUFFER CACHE



- kernel could read/write directly to/from the disk.
 - slow disk transfer rate
 - system response time is poor
- minimize the frequency of disk access by keeping a pool of data buffers
- data buffers are called as buffer cache
- it contains data in recently used blocks



POSITION OF BUFFER CACHE

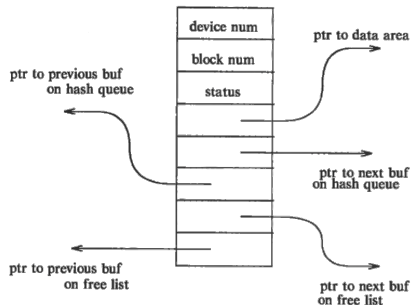


BUFFER HEADERS I



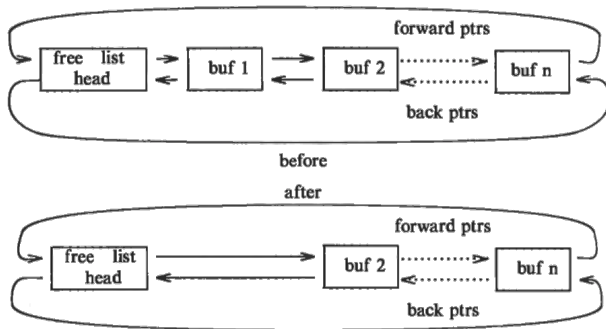
- During system initialization, the kernel allocates space for a number of buffers, configurable according to memory size and performance constraints.
- Two parts of the buffer:
 - 1 a memory array that contains data from the disk.
 - 2 “**buffer header**” that identifies the buffer.
- Data in a buffer corresponds to data in a logical disk block on a file system.
- A disk block can **never** map into more than one buffer at a time.

BUFFER HEADERS II



- **device number:** specifies logical file system
- **block number** of the data on disk
- **ptr to data area:** size must be at least as big as the size of a disk block.
- The status of a buffer is a combination of –
 - Buffer is locked / busy
 - Buffer contains valid data
 - Kernel must write the buffer contents to disk before reassigning the buffer; called as **delayed-write**
 - Kernel is currently reading or writing the contexts of the buffer to disk
 - A process is waiting for buffer to become free.
- The two set of pointers are used for traversal of the buffer queues (doubly circular linked lists).


STRUCTURE OF BUFFER POOL



Free list of buffers

- Algorithm: least recently used (LRU)
- System Boot: all buffers are on the free list
- Free buffer is taken from the head of the free list
- After use, free buffers are attached to the end of the list

STRUCTURE OF BUFFER POOL



Buffers on the Hash Queues

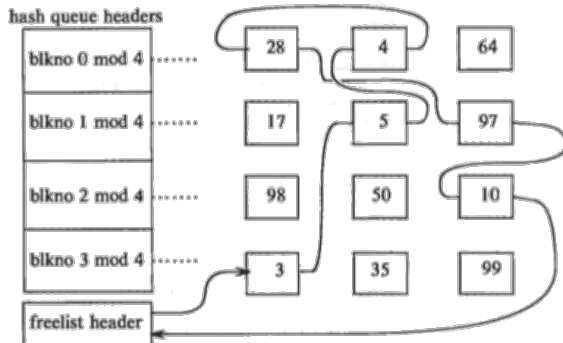
- Hash queues are doubly linked circular lists.
- Separate queues, **hashed** as a function of the device and block number
- Every disk block mapped to only one hash queue and only once
- A buffer is always on a hash queue, but it may or may not be on the free list

SCENARIOS FOR RETRIEVAL OF A BUFFER

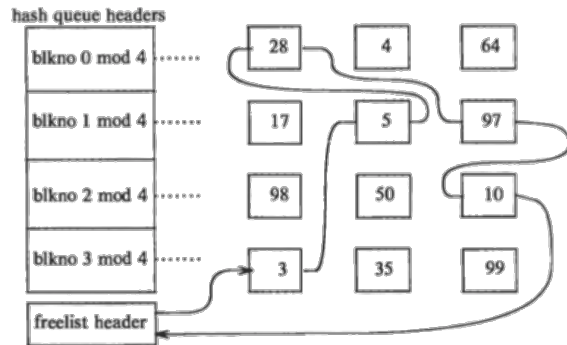


- ❶ Block is found on its hash queue and its buffer is free.
- ❷ Block could not be found on the hash queue, so a buffer from the free list is allocated.
- ❸ Block could not be found on the hash queue, and when allocating a buffer from free list, a buffer marked **delayed write** is allocated. Then the kernel must write the **delayed write** buffer to disk and allocate another buffer.
- ❹ Block could not be found on the hash queue and the free list of buffers is empty.
- ❺ Block was found on the hash queue, but its buffer is currently busy.

SCENARIO-1

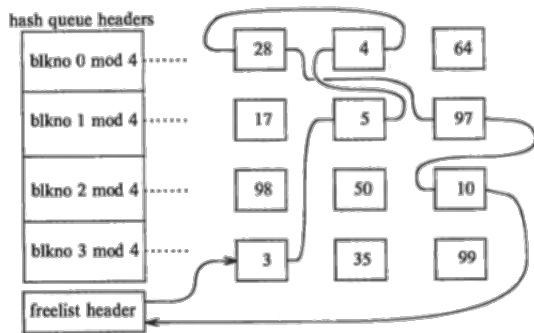


(a) Search for Block 4 on First Hash Queue

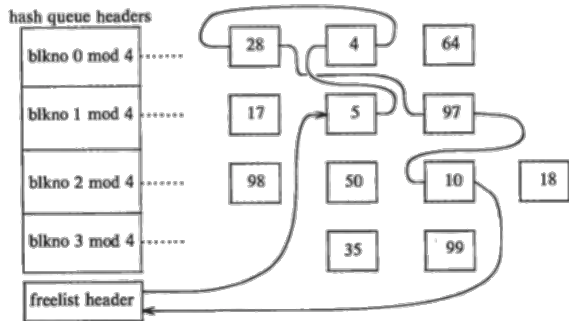


(b) Remove Block 4 from Free List

SCENARIO-2




(a) Search for Block 18 - Not in Cache




(b) Remove First Block from Free List, Assign to 18

SCENARIO-3




(a) Search for Block 18, Delayed Write Blocks on Free List



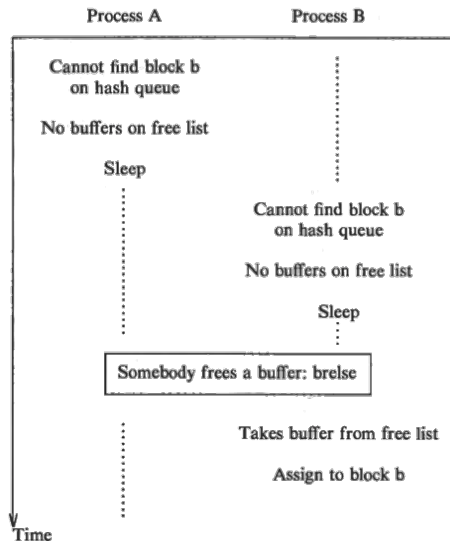
(b) Writing Blocks 3, 5, Reassign 4 to 18

SCENARIO-4




Search for Block 18, Empty Free List

SCENARIO-4 RACE FOR FREE BUFFER

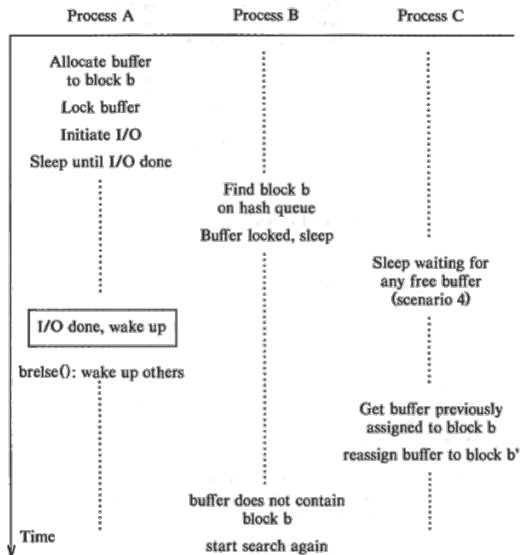


SCENARIO-5





SCENARIO-5 RACE FOR LOCKED BUFFER





Algorithm: getblk

```
while (buffer not found) {
  if (block in hash queue) {
    if (buffer busy) { // scenario 5
      sleep (event: buffer becomes free);
      continue;      // back to while loop
    }
    mark buffer busy; // scenario 1
    remove buffer from free list;
    return buffer;
  } else {
    if (there are no buffers on the free list) {
      sleep (event: any buffer becomes free);
      continue;      // back to while loop
    }
    remove buffer from free list;
    if (buffer marked for delayed write) {
      asynchronous write buffer to disk;
      continue;      // back to while loop;
    }
    remove buffer from old hash queue; // scenario 2
    put buffer onto new hash queue;
    return buffer;
  }
}
```



Algorithm: brelse

```
/*
 * Algorithm: brelse
 * Input: locked buffer
 * Output: none
 */
{
    wakeup all processes (event: waiting for any buffer to become free);
    wakeup all processes (event: waiting for this buffer to become free);

    raise processor execution level to block interrupts;

    if (buffer contents valid and buffer not old)
        enqueue buffer at end of free list;
    else
        enqueue buffer at beginning of free list;

    lower processor execution level to allow interrupts;

    unlock (buffer);
}
```



Algorithm: bread

```
/*
 * Algorithm: bread
 * Input: file system number
 *        block number
 * Output: buffer containing data
 */
{
    get buffer for block (algorithm: getblk);

    if (buffer data valid)
        return buffer;

    initiate disk read;
    sleep (event: disk read complete);

    return buffer;
}
```



Algorithm: breada

```
/* Algorithm: breada
 * Input: file system number and block number for immediate read
 *        file system number and block number for asynchronous read
 * Output: buffer containing data for immediate read
 */
{
  if (first block not in cache) {
    get buffer for first block (algorithm: bread);
    if (buffer data not valid)
      initiate disk read;
  }
  if (second block not in cache) {
    get buffer for second block (algorithm: getblk);
    if (buffer data valid)
      release buffer (algorithm: brelse);
    else
      initiate disk read;
  }
  if (first block was originally in the cache) {
    read first block (algorithm: bread);
    return buffer;
  }
  sleep (event: first buffer contains valid data);
  return buffer;
}
```



Algorithm: bwrite

```
/*
 * Algorithm: bwrite
 * Input: buffer
 * Output: none
 */
{
    initiate disk write;
    if (I/O synchronous)
    {
        sleep (event: I/O complete);
        release buffer (algorithm: brelse);
    }
    else if (buffer marked for delayed write)
        mark buffer to put at head of free list;
}
```

Thank you . . .